



Master Thesis

Virtual Tool-boxing for Robust Management of Cross-layer Heterogeneity in Complex Cyber-physical Systems


Candidate: Simone Dutto

Supervisors: prof. Stefano Di Carlo, prof. Alessandro Savino



Abstract

The final goal of this work is to analyze what can be done in terms of fault detection using hardware metrics and machine learning.



AGENDA



01

Simulation Environment

How I obtained the data

02

Data analysis and Models

How I used data to
create machine learning
models

03

OS Implementation

How I exploited the
models to create a fault
detection system

04

Conclusion and Future Works

Final thoughts and
possible extensions



01

SIMULATION ENVIRONMENT



gem5

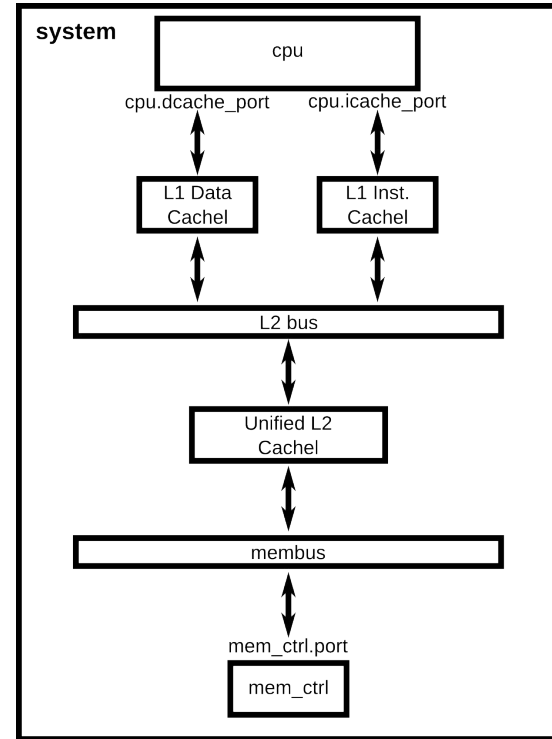
gem5 is an hardware simulator able to:

- Simulate hardware architecture
- Extract metrics from components
- Customizable to inject faults
- Run Linux Distribution



Simulated System


- Atomic Simple CPU
- x86 architecture
- L1-L2 cache
- Full System Simulation
- MiBench programs



FIMSIM

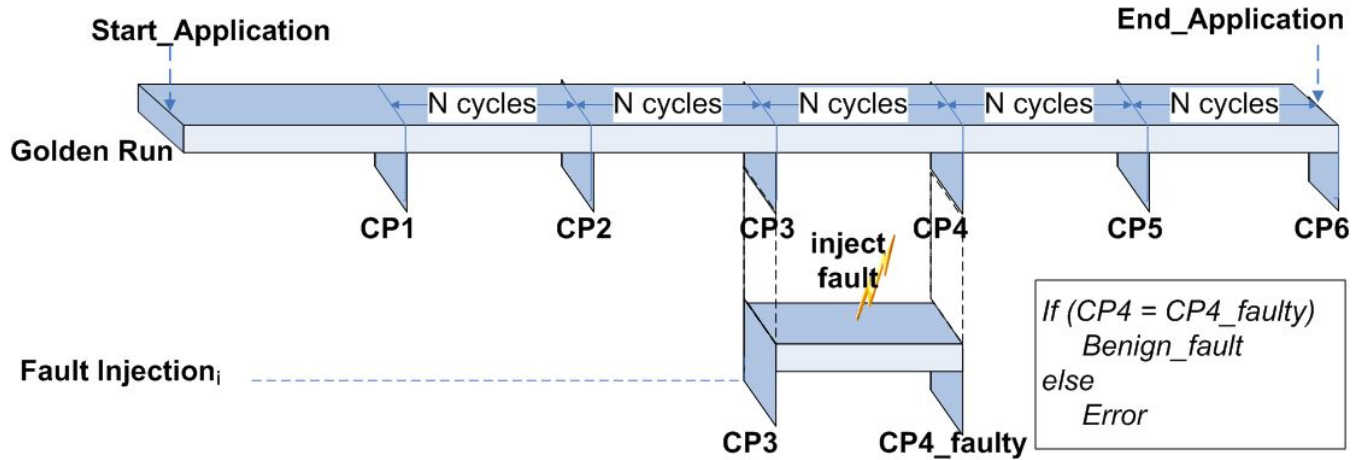
FIMSIM is a gem5 extension able to simulate the occurrence of different types of faults:

1. Stuck-at-0
2. Stuck-at-1
3. Transient



Fault Injection Results

The fault injection is evaluated comparing golden runs with faulty runs.



Data collected

- Stuck-at-0 (10k runs, 70% crash rate)
 - bitcount
 - qsort
 - basicmath
 - ssearch
- Stuck-at-1 (10k runs, 80% crash rate)
 - bitcount
 - qsort
- Transient (10k runs, 1% crash rate)
 - bitcount

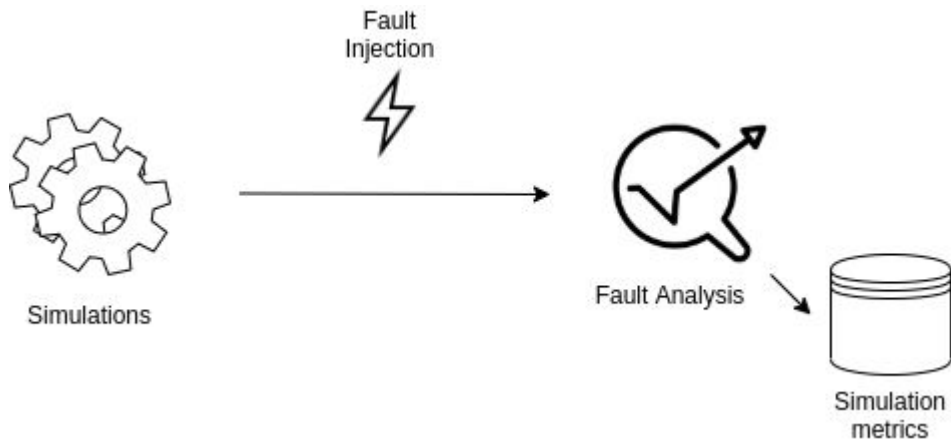


02

Data Analysis and Model



Data



```
----- Begin Simulation Statistics -----  
sim_seconds  0.047551  # Number of seconds simulated  
sim_ticks    47550712000  # Number of ticks simulated|
```

Data Preparation

Repeat for each binary and each fault

PHASE 1

Statistics files are parsed to obtain different dictionaries, which are then merged

PHASE 3

Fill NaN with column mean

PHASE 2

Drop columns with >20 NaN and entries with all NaN

PHASE 4

Remove all columns with <0.1 variance

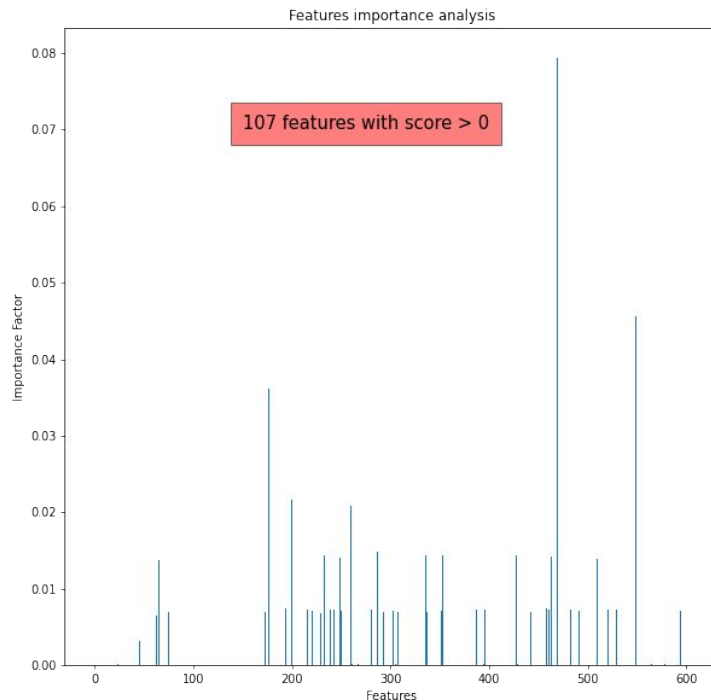


Feature Selection

Feature importance score using Random Forest Classifier.

- Avoid searching just for linear relationships
- Keep the meaning of features

I have selected the top 19 corresponding feature importance scores.



Binaries and domains

Different binaries carry different metrics distribution.

So different models and techniques are tried to cope with the differences.

| | | | |
|--|--|--|---|
| Basicmath: Simple math lib operations | QSort: Quick sort on an array | Search: Search a string inside a string | Bitcount: Count number of 1-bit in a number |
|--|--|--|---|

Model-FFNN

Simple Neural Network with 19-features selection trained on a domain for 50 epochs.

Exploit transfer learning (10 epochs) to adapt on different domains.

Very good convergence, fast to train and evaluate.

| Binary | Precision: baseline | Precision: after | Recall: baseline | Recall: after |
|-----------|---------------------|------------------|------------------|---------------|
| basicmath | 1.00 | / | 0.74 | / |
| qsort | 0.38 | 1.00 (+0.62) | 1.00 | 0.82(-0.18) |
| search | 0.46 | 0.82 (+0.36) | 1.00 | 0.33(-0.67) |
| bitcount | 1.00 | 1.00(+0.0) | 0.73 | 0.74(+0.01) |

Model-SSAE

Sparse Stacked AutoEncoder, based on CNN and full features selection.

There the auto-encoder is trained to with unlabelled data on non-faulty runs.

The principal concept is to have the encoder able to distantiate faulty and non-faulty runs the best it can before training the classifier on actual labelled data.

| Binary | Precision: before | Precision: after | Recall: before | Recall: after |
|-----------|-------------------|------------------|----------------|---------------|
| basicmath | 1.00 | \ | 0.80 | \ |
| qsort | 1.00 | 1.00(+0.0) | 0.20 | 0.78 |
| search | 0.86 | 0.83(-0.03) | 0.11 | 0.32 |
| bitcount | 0.12 | 1.00(+0.88) | 0.23 | 0.77 |

Model-DANN


Domain Adversarial Neural Network, based on CNN and full features selection.

The idea is to have source and target domains, the first entries are labelled, the second are not.

The catch here is to train the feature extractor to perform on the task without considering the domain using an adversarial approach.

| Binary | Precision: before | Precision: after | Recall: before | Recall: after |
|-----------|-------------------|------------------|----------------|---------------|
| basicmath | 1 | \ | 0.80 | \ |
| qsort | 0.23 | 1(+0.77) | 1 | 0.77 |
| search | 0.36 | 0.46(+0.10) | 0.8 | 1 |
| bitcount | 0.33 | 0.38(+0.05) | 0.8 | 1 |

Results





03

OS Implementation



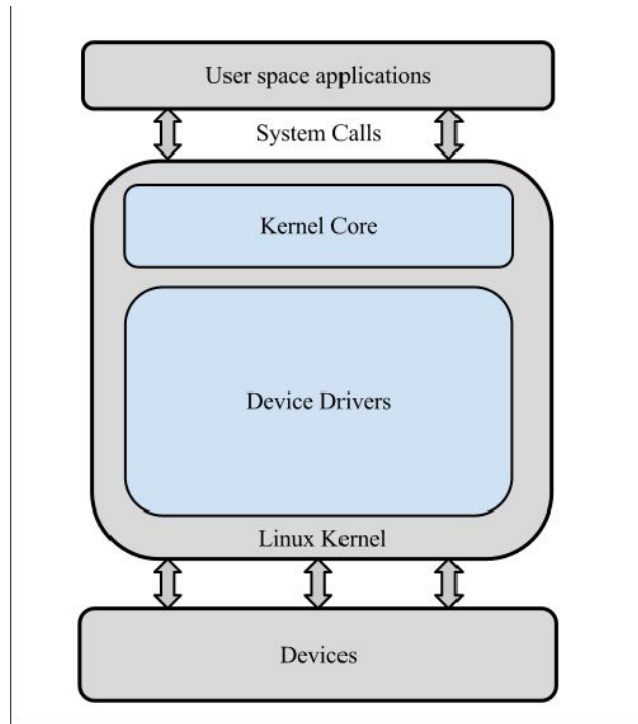
Kernel Module

Fast as the functions were implemented in the kernel.

The change is not bounded to the kernel.

Loaded only if needed.

Easy interactions with user-level applications.



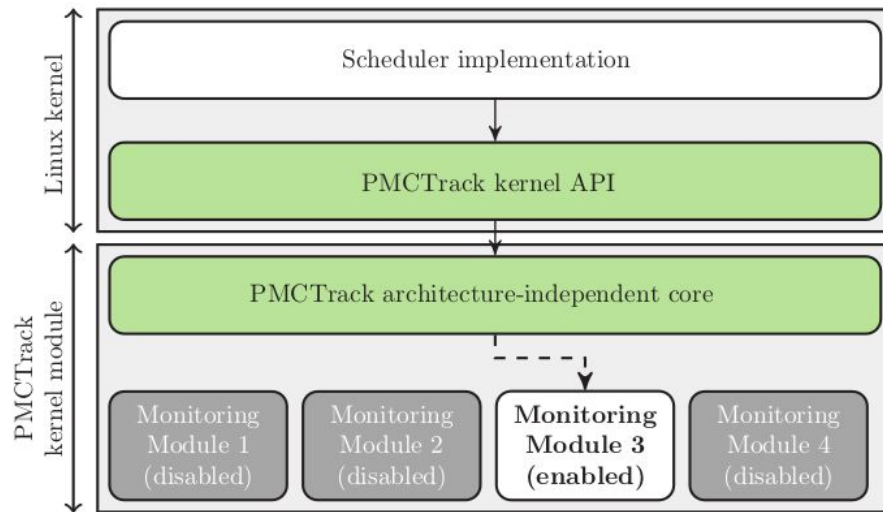
PMCTrack

Kernel module to extract Performance Monitor Counters.

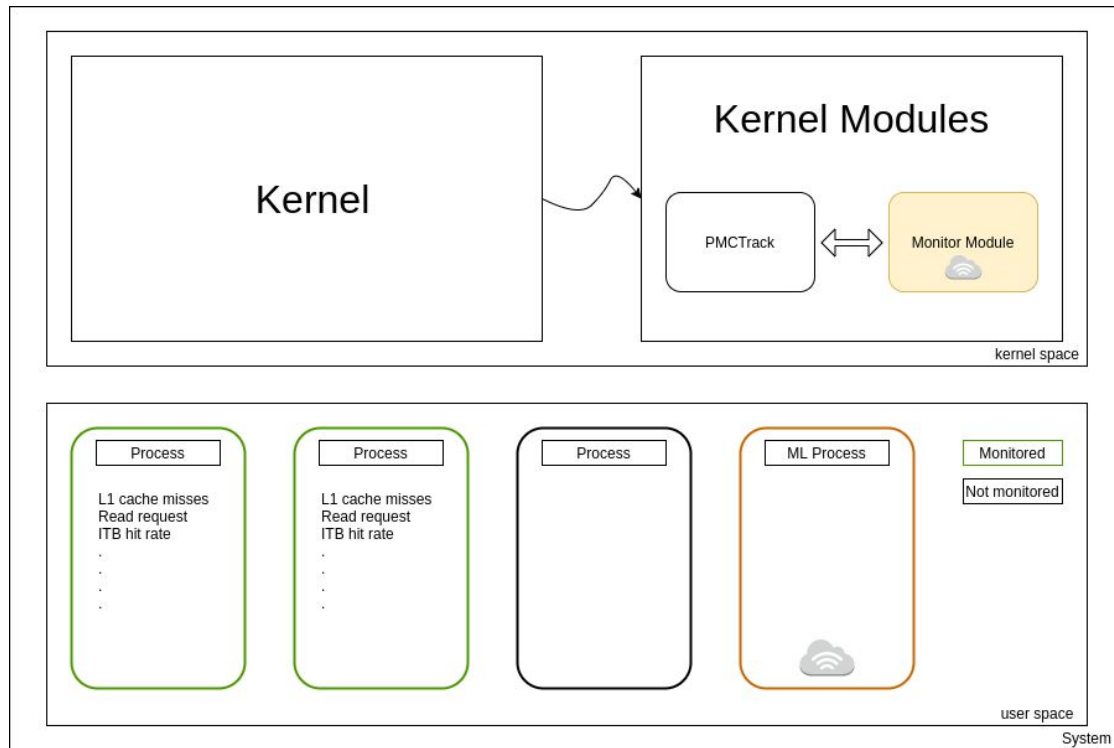
Minor changes to the kernel, all implemented at kernel module level.

Expose metrics to kernel and user-level applications.

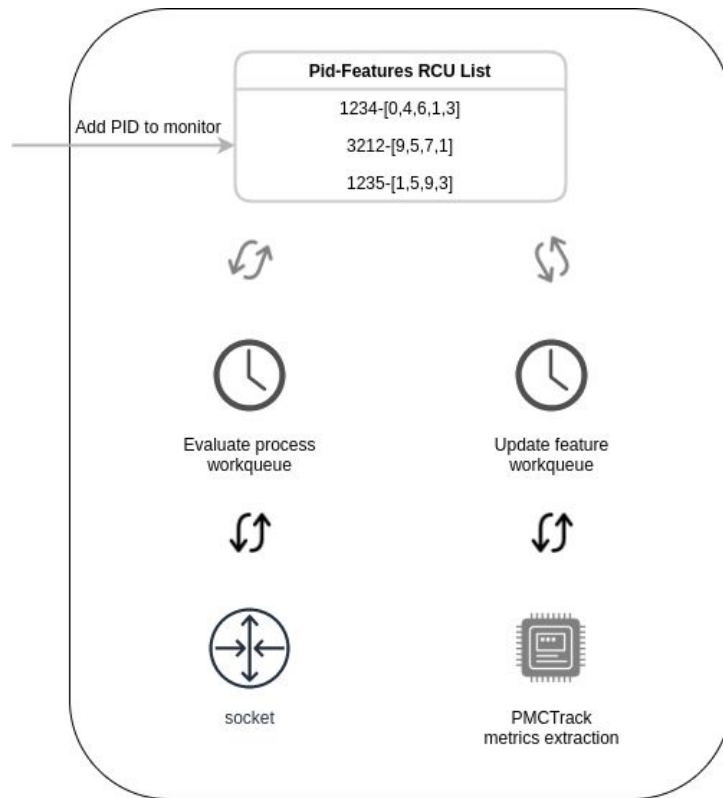
Process-wise.



Monitor Framework



Monitor Module





04

Conclusion and Future works



Conclusion

It is definitely possible to state the fault detection is doable though hardware metrics.

Even a simple model works very well with a 19-features selection.

The OS implementation needs other metrics implemented to be working.

In addition, a lot can be done to improve the faults detection framework.



Future works

Collect more data with different binaries and fault types.

Modify simulation output to resemble stream output.

Build time-based models to work with stream data.

Add metrics support to PMCTrack.

Incapsulate the model evaluation kernel-side.

Performance analysis with monitor framework.

The work on machine learning models has been submitted to a conference.

Contacts

To have more details on the whole project, it is possible to:

- Consult my thesis
- Have a look at my GitHub page with all the code related at github.com/SimoneDutto
- Write me an email simonedutto8@gmail.com